

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow [@Openwall](#) on Twitter for new release announcements and other news

[<prev](#)] [\[next>](#)] [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Tue, 14 Jun 2022 13:05:08 +1000  
From: Michael Ellerman <mpe@...erman.id.au>  
To: oss-security@...ts.openwall.com  
Subject: CVE-2022-32981: Linux kernel for powerpc 32-bit, buffer overflow in  
ptrace PEEKUSER/POKEUSER

The Linux kernel for powerpc 32-bit has a buffer overflow in the handling of ptrace  
PEEKUSER/POKEUSER when accessing floating point registers.

The fix for mainline is:  
<https://git.kernel.org/pub/scm/linux/kernel/git/powerpc/linux.git/commit/?id=8e127844446fc97778a5e5c99bcalce0bbc5ec9>

Which is included in v5.19-rc2.

Stable backports have been posted.

A test case is included below, it will report if the system is correctly  
patched, it is safe to run on an unpatched system.

cheers

---

```
#undef NDEBUG
#include <assert.h>
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/ptrace.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <sys/syscall.h>
#include <unistd.h>
#include <sys/ipc.h>
#include <sys/shm.h>

static double expected = 0.123456L;

static int child(int shm_id)
{
    int *cptr = shmat(shm_id, NULL, 0);

    asm volatile (
        "lfd    %%f0, 0(%0)    ;"
        "lfd    %%f1, 0(%0)    ;"
        "li     %%r9, 1        ;"
        "stw    %%r9, 0(%1)    ;"
        "1:"
        "lwz    %%r9, 0(%2)    ;"
        "cmpwi  %%r9, 0        ;"
        "beq    1b             ;"
        : // outputs
        : // inputs
        "b" (&expected), "b" (&cptr[1]), "b" (&cptr[0])
        : // clobbers
        "memory", "r9", "fr0", "fr1"
    );

    return 0;
}

int start_trace(pid_t child)
{
    int ret;
```

```

    ret = ptrace(PTRACE_ATTACH, child, NULL, NULL);
    if (ret) {
        perror("ptrace(PTRACE_ATTACH) failed");
        return -1;
    }
    ret = waitpid(child, NULL, 0);
    if (ret != child) {
        perror("waitpid() failed");
        return -1;
    }
    return 0;
}

int stop_trace(pid_t child)
{
    int ret;

    ret = ptrace(PTRACE_DETACH, child, NULL, NULL);
    if (ret) {
        perror("ptrace(PTRACE_DETACH) failed");
        return -1;
    }
    return 0;
}

long raw_ptrace(enum __ptrace_request request, pid_t pid, unsigned long addr, void *data)
{
    return syscall(__NR_ptrace, request, pid, (void *)addr, data);
}

#define PEEKS_PER_FPR    (sizeof(__u64) / sizeof(unsigned long))

int peek_fpr(pid_t child, int frnum, __u64 *fpr)
{
    unsigned long *p, addr;
    int i, fpindex;
    long ret;

    fpindex = PEEKS_PER_FPR * frnum;

    p = (unsigned long *)fpr;
    for (i = 0; i < PEEKS_PER_FPR; i++, p++) {
        addr = sizeof(unsigned long) * (PT_FPR0 + fpindex + i);
        ret = raw_ptrace(PTRACE_PEEKUSER, child, addr, p);
        if (ret) {
            perror("ptrace(PTRACE_PEEKUSR) failed");
            return -1;
        }
    }

    return 0;
}

int parent(pid_t child)
{
    double f0, f1;

    assert(start_trace(child) == 0);

    assert(peek_fpr(child, 0, (__u64 *)&f0) == 0);
    assert(peek_fpr(child, 1, (__u64 *)&f1) == 0);

    assert(stop_trace(child) == 0);

    printf("expected = %e\n", f0);
    printf("f0          = %e\n", f0);
    printf("f1          = %e\n", f1);

    if (f0 != expected || f1 != expected) {
        printf("FAIL - values don't match! Kernel is buggy.\n");
        return -1;
    }

    printf("OK - values match\n");
}

```

```

        return 0;
    }

int main(void)
{
    int shm_id, ret, status, *pptr;
    pid_t pid;

    shm_id = shmget(IPC_PRIVATE, sizeof(int) * 2, 0777|IPC_CREAT);
    assert(shm_id != -1);

    pid = fork();
    assert(pid >= 0);

    if (pid == 0)
        exit(child(shm_id));

    pptr = shmat(shm_id, NULL, 0);

    // Wait for child to signal us to continue
    while (!pptr[1])
        asm volatile("" : : : "memory");

    ret = parent(pid);
    if (ret) {
        kill(pid, SIGTERM);
        shmdt((void *)pptr);
        shmctl(shm_id, IPC_RMID, NULL);
        return -1;
    }

    // Signal child to exit
    pptr[0] = 1;
    shmdt((void *)pptr);

    ret = wait(&status);
    shmctl(shm_id, IPC_RMID, NULL);

    assert(ret != -1 && WIFEXITED(status) && !WEXITSTATUS(status));

    return 0;
}

```

Powered by [blists](#) - more mailing lists

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

