



index : kernel/git/torvalds/linux.git

Linux kernel source tree

master switch
Linus Torvalds

about summary refs log tree commit diff stats

log msg search

author Kazuki Takiguchi <takiguchi.kazuki171@gmail.com> 2022-11-23 14:36:00 -0500
committer Paolo Bonzini <pbonzini@redhat.com> 2022-11-23 18:50:08 -0500
commit 47b0c2e4c220f2251fd8dcfbb44479819c715e15 (patch)
tree 04a4bb3b263c55a437acbbd373116db0cc24addb
parent d79b483193c276f9b6863f69735b97de12ced621 (diff)
download linux-47b0c2e4c220f2251fd8dcfbb44479819c715e15.tar.gz

diff options

context: 3

space: include

mode: unified

KVM: x86/mmu: Fix race condition in direct_page_fault

make_mmu_pages_available() must be called with mmu_lock held for write. However, if the TDP MMU is used, it will be called with mmu_lock held for read. This function does nothing unless shadow pages are used, so there is no race unless nested TDP is used. Since nested TDP uses shadow pages, old shadow pages may be zapped by this function even when the TDP MMU is enabled. Since shadow pages are never allocated by kvm_tdp_mmu_map(), a race condition can be avoided by not calling make_mmu_pages_available() if the TDP MMU is currently in use.

I encountered this when repeatedly starting and stopping nested VM. It can be artificially caused by allocating a large number of nested TDP SPTEs.

For example, the following BUG and general protection fault are caused in the host kernel.

```
pte_list_remove: 00000000cd54fc10 many->many
-----[ cut here ]-----
kernel BUG at arch/x86/kvm/mmu/mmu.c:963!
invalid opcode: 0000 [#1] PREEMPT SMP NOPTI
RIP: 0010:pte_list_remove.cold+0x16/0x48 [kvm]
Call Trace:
<TASK>
drop_spte+0xe0/0x180 [kvm]
mmu_page_zap_pte+0x4f/0x140 [kvm]
__kvm_mmu_prepare_zap_page+0x62/0x3e0 [kvm]
kvm_mmu_zap_oldest_mmu_pages+0x7d/0xf0 [kvm]
direct_page_fault+0x3cb/0x9b0 [kvm]
kvm_tdp_page_fault+0x2c/0xa0 [kvm]
kvm_mmu_page_fault+0x207/0x930 [kvm]
npf_interception+0x47/0xb0 [kvm_amd]
svm_invoke_exit_handler+0x13c/0x1a0 [kvm_amd]
svm_handle_exit+0xfc/0x2c0 [kvm_amd]
kvm_arch_vcpu_ioctl_run+0xa79/0x1780 [kvm]
kvm_vcpu_ioctl+0x29b/0x6f0 [kvm]
__x64_sys_ioctl+0x95/0xd0
do_syscall_64+0x5c/0x90
```

```
general protection fault, probably for non-canonical address
0xdead000000000122: 0000 [#1] PREEMPT SMP NOPTI
RIP: 0010:kvm_mmu_commit_zap_page.part.0+0x4b/0xe0 [kvm]
Call Trace:
<TASK>
kvm_mmu_zap_oldest_mmu_pages+0xae/0xf0 [kvm]
direct_page_fault+0x3cb/0x9b0 [kvm]
kvm_tdp_page_fault+0x2c/0xa0 [kvm]
kvm_mmu_page_fault+0x207/0x930 [kvm]
npf_interception+0x47/0xb0 [kvm_amd]
```

CVE: CVE-2022-45869
Fixes: a2855afc7ee8 ("KVM: x86/mmu: Allow parallel page faults for the TDP MMU")
Signed-off-by: Kazuki Takiguchi <takiguchi.kazuki171@gmail.com>
Cc: stable@vger.kernel.org
Signed-off-by: Paolo Bonzini <pbonzini@redhat.com>

Diffstat

-rw-r--r- arch/x86/kvm/mmu/mmu.c 13

1 files changed, 7 insertions, 6 deletions

```
diff --git a/arch/x86/kvm/mmu/mmu.c b/arch/x86/kvm/mmu/mmu.c
index 1ccb769f62af3..b6f96d47e596d 100644
--- a/arch/x86/kvm/mmu/mmu.c
+++ b/arch/x86/kvm/mmu/mmu.c
@@ -2443,6 +2443,7 @@ static bool __kvm_mmu_prepare_zap_page(struct kvm *kvm,
 {
     bool list_unstable, zapped_root = false;

+    lockdep_assert_held_write(&kvm->mmu_lock);
     trace_kvm_mmu_prepare_zap_page(sp);
     ++kvm->stat.mmu_shadow_zapped;
     *nr_zapped = mmu_zap_unsync_children(kvm, sp, invalid_list);
@@ -4262,14 +4263,14 @@ static int direct_page_fault(struct kvm_vcpu *vcpu, struct kvm_page_fault *fault
     if (is_page_fault_stale(vcpu, fault, mmu_seq))
         goto out_unlock;

-    r = make_mmu_pages_available(vcpu);
-    if (r)
-        goto out_unlock;
-
     if (is_tdp_mmu_fault)
         if (is_tdp_mmu_fault) {
             r = kvm_tdp_mmu_map(vcpu, fault);
         } else {
             r = make_mmu_pages_available(vcpu);
             if (r)
                 goto out_unlock;
             r = __direct_map(vcpu, fault);
         }

     out_unlock:
     if (is_tdp_mmu_fault)
```