

linux-media.vger.kernel.org archive mirror

search help / color / mirror / Atom feed

From: imv4bel@gmail.com
To: mchehab@kernel.org
Cc: Hyunwoo Kim <imv4bel@gmail.com>,
kernel@tuxforce.de, linux-media@vger.kernel.org,
linux-usb@vger.kernel.org, cai.huojing@linux.dev, tiwai@suse.de
Subject: [PATCH 3/4] media: dvb-core: Fix use-after-free due to race condition occurring in dvb_register_device()
Date: Tue, 15 Nov 2022 05:18:21 -0800 [thread overview]
Message-ID: <20221115131822.6640-4-imv4bel@gmail.com> (raw)
In-Reply-To: <20221115131822.6640-1-imv4bel@gmail.com>

From: Hyunwoo Kim <imv4bel@gmail.com>

dvb_register_device() dynamically allocates fops with kmemdup()
to set the fops->owner
And these fops are registered in 'file->f_ops' using replace_fops()
in the dvb_device_open() process, and kfree()d in dvb_free_device().

However, it is not common to use dynamically allocated fops instead
of 'static const' fops as an argument of replace_fops(),
and UAF may occur.
These UAFs can occur on any dvb type using dvb_register_device(),
such as dvb_dvr, dvb_demux, dvb_frontend, dvb_net, etc.

So, instead of kfree() the fops dynamically allocated in
dvb_register_device() in dvb_free_device() called during the
.disconnect() process, kfree() it collectively in exit_dvbdev()
called when the dvbdev.c module is removed.

Signed-off-by: Hyunwoo Kim <imv4bel@gmail.com>

```
drivers/media/dvb-core/dvbdev.c | 83 ++++++-----
include/media/dvbdev.h          | 15 +++++
2 files changed, 77 insertions(+), 21 deletions(-)

diff --git a/drivers/media/dvb-core/dvbdev.c b/drivers/media/dvb-core/dvbdev.c
index 675d877a67b2..424cf92c068e 100644
--- a/drivers/media/dvb-core/dvbdev.c
+++ b/drivers/media/dvb-core/dvbdev.c
@@ -27,6 +27,7 @@
#include <media/tuner.h>

static DEFINE_MUTEX(dvbdev_mutex);
+static LIST_HEAD(dvbdevfops_list);
static int dvbdev_debug;

module_param(dvbdev_debug, int, 0644);
@@ -448,14 +449,15 @@ int dvb_register_device(struct dvb_adapter *adap, struct dvb_device **pdvbdev,
enum dvb_device_type type, int demux_sink_pads)
{
    struct dvb_device *dvbdev;
-    struct file_operations *dvbdevfops;
+    struct file_operations *dvbdevfops = NULL;
+    struct dvbdevfops_node *node, *new_node;
    struct device *clsdev;
    int minor;
    int id, ret;

    mutex_lock(&dvbdev_register_lock);

    if ((id = dvbdev_get_free_id(adap, type)) < 0) {
+    if ((id = dvbdev_get_free_id(adap, type)) < 0) {
        mutex_unlock(&dvbdev_register_lock);
        *pdvbdev = NULL;
        pr_err("%s: couldn't find free device id\n", func);
@@ -463,18 +465,19 @@ }

    *pdvbdev = dvbdev = kzalloc(sizeof(*dvbdev), GFP_KERNEL);

    if (!dvbdev) {
        mutex_unlock(&dvbdev_register_lock);
        return -ENOMEM;
    }

    dvbdevfops = kmemdup(template->fops, sizeof(*dvbdevfops), GFP_KERNEL);
+    /*
+     * When a device of the same type is probe()d more than once,
+     * the first allocated fops are used. This prevents memory leaks
+     * that can occur when the same device is probe()d repeatedly.
+     */
    list_for_each_entry(node, &dvbdevfops_list, list_head) {
        if (node->fops->owner == adap->module &&
            node->type == type &&
            node->template == template) {
            dvbdevfops = node->fops;
            break;
        }
    }

    if (!dvbdevfops) {
        kfree(dvbdev);
        mutex_unlock(&dvbdev_register_lock);
        return -ENOMEM;
    }
    if (dvbdevfops == NULL) {
        dvbdevfops = kmemdup(template->fops, sizeof(*dvbdevfops), GFP_KERNEL);
        if (!dvbdevfops) {
            kfree(dvbdev);
            mutex_unlock(&dvbdev_register_lock);
            return -ENOMEM;
        }
        new_node = kzalloc(sizeof(struct dvbdevfops_node), GFP_KERNEL);
        if (!new_node) {
            kfree(dvbdevfops);
            kfree(dvbdev);
            mutex_unlock(&dvbdev_register_lock);
            return -ENOMEM;
        }
        new_node->fops = dvbdevfops;
        new_node->type = type;
        new_node->template = template;
        list_add_tail(&new_node->list_head, &dvbdevfops_list);
    }

    memcpy(dvbdev, template, sizeof(struct dvb_device));
@@ -484,20 +484,20 @@ int dvb_register_device(struct dvb_adapter *adap, struct dvb_device **pdvbdev,
dvbdev->priv = NULL;
dvbdev->fops = dvbdevfops;
init_waitqueue_head(&dvbdev->wait_queue);

    dvbdevfops->owner = adap->module;

    list_add_tail(&dvbdev->list_head, &adap->device_list);

    down_write(&minor_rwsem);
#ifdef CONFIG_DVB_DYNAMIC_MINORS
    for (minor = 0; minor < MAX_DVB_MINORS; minor++)
        if (dvb_minors[minor] == NULL)
            break;
#endif

    if (minor == MAX_DVB_MINORS) {
        if (new_node) {
            list_del(&new_node->list_head);
            kfree(dvbdevfops);
            kfree(new_node);
        }
    }
}
```

```

+         }
+         list_del (&dvbdev->list_head);
-         kfree (dvbdevfops);
+         kfree (dvbdev);
+         up_write (&minor_rwlock);
+         mutex_unlock (&dvbdev_register_lock);
@@ -506,41 +535,46 @@ int dvb_register_device (struct dvb_adapter *adap, struct dvb_device **pdvbdev,
+         #else
+         minor = nums2minor (adap->num, type, id);
+         #endif
-         dvbdev->minor = minor;
+         dvb_minors[minor] = dvbdev;
+         up_write (&minor_rwlock);
-         ret = dvb_register_media_device (dvbdev, type, minor, demux_sink_pads);
+         if (ret) {
+             pr_err ("%s: dvb_register_media_device failed to create the mediagraph\n",
+                     __func__);
-         if (new_node) {
+             if (new_node) {
+                 list_del (&new_node->list_head);
+                 kfree (dvbdevfops);
+                 kfree (new_node);
+             }
+             dvb_media_device_free (dvbdev);
+             list_del (&dvbdev->list_head);
-             kfree (dvbdevfops);
+             kfree (dvbdev);
+             mutex_unlock (&dvbdev_register_lock);
+             return ret;
+         }
-         mutex_unlock (&dvbdev_register_lock);
-         clsdev = device_create (dvb_class, adap->device,
+             MKDEV (DVB_MAJOR, minor),
+             dvbdev, "dvb%d.%s%d", adap->num, dnames[type], id);
+         if (IS_ERR (clsdev)) {
+             pr_err ("%s: failed to create device dvb%d.%s%d (%ld)\n",
+                     __func__, adap->num, dnames[type], id, PTR_ERR (clsdev));
+             if (new_node) {
+                 list_del (&new_node->list_head);
+                 kfree (dvbdevfops);
+                 kfree (new_node);
+             }
+             dvb_media_device_free (dvbdev);
+             list_del (&dvbdev->list_head);
-             kfree (dvbdevfops);
+             kfree (dvbdev);
+             return PTR_ERR (clsdev);
+         }
+         dprintk ("DVB: register adapter%d/%s%d @ minor: %i (0x%02x)\n",
+                 adap->num, dnames[type], id, minor, minor);
+         mutex_unlock (&dvbdev_register_lock);
+         return 0;
+     }
+ EXPORT_SYMBOL (dvb_register_device);
@@ -569,7 +603,6 @@ void dvb_free_device (struct dvb_device *dvbdev)
+         if (!dvbdev)
+             return;
-         kfree (dvbdev->fops);
+         kfree (dvbdev);
+     }
+ EXPORT_SYMBOL (dvb_free_device);
@@ -1061,9 +1094,17 @@ static int __init init_dvbdev (void)
+         static void __exit exit_dvbdev (void)
+         {
+             struct dvbdevfops_node *node, *next;
+             class_destroy (dvb_class);
+             cdev_del (&dvb_device_cdev);
+             unregister_chrdev_region (MKDEV (DVB_MAJOR, 0), MAX_DVB_MINORS);
+             list_for_each_entry_safe (node, next, &dvbdevfops_list, list_head) {
+                 list_del (&node->list_head);
+                 kfree (node->fops);
+                 kfree (node);
+             }
+         }
+         subsys_initcall (init_dvbdev);
+         diff --git a/include/media/dvbdev.h b/include/media/dvbdev.h
+         index 2f6b0861322a..1e5413303705 100644
+         --- a/include/media/dvbdev.h
+         +++ b/include/media/dvbdev.h
+         @@ -187,6 +187,21 @@ struct dvb_device {
+             void *priv;
+         };
+         /**
+         * struct dvbdevfops_node - fops nodes registered in dvbdevfops_list
+         *
+         * @fops: Dynamically allocated fops for ->owner registration
+         * @type: type of dvb device
+         * @template: dvb device used for registration
+         * @list_head: list_head for dvbdevfops_list
+         */
+         struct dvbdevfops_node {
+             struct file_operations *fops;
+             enum dvb_device_type type;
+             const struct dvb_device *template;
+             struct list_head list_head;
+         };
+         /**
+         * dvb_register_adapter - Registers a new DVB adapter
+         */
+         --
+         2.25.1

```

next prev parent reply other threads: [~2022-11-15 13:19 UTC|newest]

Thread overview: 6+ messages / expand[flat|nested] mbox.gz Atom feed top

2022-11-15 13:18 [PATCH 0/4] Fix multiple race condition vulnerabilities in dvb-core and device driver imv4bel
2022-11-15 13:18 ` [PATCH 1/4] media: dvb-core: Fix use-after-free due to race condition occurring in dvb_frontend imv4bel
2022-11-15 13:18 ` [PATCH 2/4] media: dvb-core: Fix use-after-free due to race condition occurring in dvb_net imv4bel
2022-11-15 13:18 ` imv4bel [this message]
2022-11-17 4:16 ` [PATCH 3/4] media: dvb-core: Fix use-after-free due to race condition occurring in dvb_register_device() Dan Carpenter
2022-11-15 13:18 ` [PATCH 4/4] media: ttusb-dec: Fix memory leak in ttusb_dec_exit_dvb() imv4bel

find likely ancestor, descendant, or conflicting patches for this message:

dfblob:675d877a67b dfblob:2f6b0861322 dfblob:424cf92c068
dfblob:1e541330370

(help)

Reply instructions:

You may reply publicly to [this message](#) via plain-text email using any one of the following methods:

* Save the following mbox file, import it into your mail client,
and reply-to-all from there: [mbox](#)

Avoid top-posting and favor interleaved quoting:
https://en.wikipedia.org/wiki/Posting_style#Interleaved_style

* Reply using the **--to**, **--cc**, and **--in-reply-to**
switches of git-send-email(1):

```
git send-email \
  --in-reply-to=20221115131822.6640-4-imv4bel@gmail.com \
  --to=imv4bel@gmail.com \
  --cc=cai.huqing@linux.dev \
  --cc=kernel@tuxforce.de \
  --cc=linux-media@vger.kernel.org \
  --cc=linux-usb@vger.kernel.org \
  --cc=mchehab@kernel.org \
  --cc=tiwai@suse.de \
  /path/to/YOUR_REPLY
```

<https://kernel.org/pub/software/scm/git/docs/git-send-email.html>

* If your mail client supports setting the **In-Reply-To** header
via mailto: links, try the [mailto: link](#)

Be sure your reply has a **Subject:** header at the top and a blank line before the message body.

This is a public inbox, see [mirroring instructions](#)
for how to clone and mirror all data and code used for this inbox;
as well as URLs for NNTP newsgroup(s).