

From: Hyunwoo Kim <imv4bel@gmail.com>
 To: eli.billauer@gmail.com, arnd@arndb.de, gregkh@linuxfoundation.org
 Cc: linux-kernel@vger.kernel.org
 Subject: [PATCH] char: xillybus: Fix use-after-free in xillyusb_open()
 Date: Sat, 22 Oct 2022 10:54:04 +0700 [thread overview]
 Message-ID: <20221022175404.GA375335@ubuntu> (raw)

A race condition may occur if the user physically removes the USB device while calling open() for this device node.

This is a race condition between the xillyusb_open() function and the xillyusb_disconnect() function, which may eventually result in UAF.

So, add a mutex to the xillyusb_open() and xillyusb_disconnect() functions to avoid race condition.

Signed-off-by: Hyunwoo Kim <imv4bel@gmail.com>

```
---
drivers/char/xillybus/xillyusb.c | 21 ++++++
1 file changed, 20 insertions(+), 1 deletion(-)

diff --git a/drivers/char/xillybus/xillyusb.c b/drivers/char/xillybus/xillyusb.c
index 39bcbfd908b4..d615f74dcf36 100644
--- a/drivers/char/xillybus/xillyusb.c
+++ b/drivers/char/xillybus/xillyusb.c
@@ -63,6 +63,7 @@ static const struct usb_device_id xillyusb_table[] = {
 };

MODULE_DEVICE_TABLE(usb, xillyusb_table);
+static DEFINE_MUTEX(disconnect_mutex);

struct xillyusb_dev;

@@ -1237,9 +1238,13 @@ static int xillyusb_open(struct inode *inode, struct file *filp)
{
    int rc;
    int index;

+    mutex_lock(&disconnect_mutex);
+
    rc = xillybus_find_inode(inode, (void **)&xdev, &index);
-    if (rc)
+    if (rc) {
+        mutex_unlock(&disconnect_mutex);
+        return rc;
+    }

    chan = xdev->channels[index];
    filp->private_data = chan;
@@ -1379,6 +1384,8 @@ static int xillyusb_open(struct inode *inode, struct file *filp)
    request_read_anything(chan, OPCODE_SET_PUSH);
}

+    mutex_unlock(&disconnect_mutex);
+
    return 0;

unfifo:
@@ -1410,10 +1417,14 @@ static int xillyusb_open(struct inode *inode, struct file *filp)
    kref_put(&xdev->kref, cleanup_dev);

+    mutex_unlock(&disconnect_mutex);
+
    return rc;

unmutex_fail:
    mutex_unlock(&chan->lock);
    mutex_unlock(&disconnect_mutex);
+
    return rc;
}

@@ -1698,6 +1709,8 @@ static int xillyusb_release(struct inode *inode, struct file *filp)
struct xillyusb_dev *xdev = chan->xdev;
int rc_read = 0, rc_write = 0;

+    mutex_lock(&disconnect_mutex);
+
    if (filp->f_mode & FMODE_READ) {
        struct xillyfifo *in_fifo = chan->in_fifo;

@@ -1760,6 +1773,8 @@ static int xillyusb_release(struct inode *inode, struct file *filp)
    kref_put(&xdev->kref, cleanup_dev);

+    mutex_unlock(&disconnect_mutex);
+
    return rc_read ? rc_read : rc_write;
}

@@ -2172,6 +2187,8 @@ static void xillyusb_disconnect(struct usb_interface *interface)
{
    int rc;
    int i;

+    mutex_lock(&disconnect_mutex);
+
    xillybus_cleanup_chrdev(xdev, &interface->dev);
}

/*
@@ -2228,6 +2245,8 @@ static void xillyusb_disconnect(struct usb_interface *interface)
xdev->dev = NULL;

    kref_put(&xdev->kref, cleanup_dev);
+
+    mutex_unlock(&disconnect_mutex);
+
}

static struct usb_driver xillyusb_driver = {
--
2.25.1
```

Dear,

This race condition can occur in the flow of:

```
...
    cpu0                                cpu1
1. xillyusb_open()
   xillybus_find_inode()
   mutex_lock(&unit_mutex);
   unit = iter;
   mutex_unlock(&unit_mutex);

2. xillyusb_disconnect()
   xillybus_cleanup_chrdev()
   mutex_lock(&unit_mutex);
   kfree(unit);
   mutex_unlock(&unit_mutex);

3. *private_data = unit->private_data; // UAF

...
```

Because the interval between 1 and 3 in xillyusb_open() is very short, this race condition is usually rare.

However, if someone wants to trigger this UAF, they can use the technique presented in this paper:

<https://www.usenix.org/system/files/sec21-lee-yoochan.pdf>

This is a method to increase the execution time between No. 1 and No. 3 using the Reschedule IPI.
This means you will be able to trigger the UAF much more reliably.

Here is the KASAN log:

```
...
[ 66.645661] =====
[ 66.645678] BUG: KASAN: use-after-free in xillybus_find_inode+0x2c6/0x2ea [xillybus_class]
[ 66.645712] Read of size 8 at addr ffff881be45f90 by Task xillybusb_test/2143

[ 66.645735] CPU: 2 PID: 2143 Comm: xillybusb_test Not tainted 6.1.0-rc1+ #10
[ 66.645752] Hardware name: Gigabyte Technology Co., Ltd. B460MDS3H/B460M DS3H, BIOS F3 05/27/2020
[ 66.645762] Call Trace:
[ 66.645772] <TASK>
[ 66.645783] dump_stack_lvl+0x49/0x63
[ 66.645808] print_report+0x177/0x46e
[ 66.645828] ? kasan_complete_mode_report_info+0x7c/0x210
[ 66.645846] ? xillybus_find_inode+0x2c6/0x2ea [xillybus_class]
[ 66.645871] kasan_report+0xb0/0x140
[ 66.645891] ? xillybus_find_inode+0x2c6/0x2ea [xillybus_class]
[ 66.645917] ? _asan_report_load8_noabort+0x14/0x20
[ 66.645932] xillybus_find_inode+0x2c6/0x2ea [xillybus_class]
[ 66.645960] xillyusb_open+0xa6/0x1030 [xillyusb]
[ 66.645990] ? endpoint_alloc+0x6d0/0x6d0 [xillyusb]
[ 66.646014] ? kasan_check_write+0x14/0x20
[ 66.646028] ? raw_spin_lock+0x88/0xe0
[ 66.646044] ? try_module_get.part.0+0xb8/0x1a0
[ 66.646062] chrdev_open+0x230/0x6d0
[ 66.646082] ? cdev_device_add+0x1f0/0x1f0
[ 66.646099] ? fsnotify_perm.part.0+0x1d9/0x4e0
[ 66.646118] do_dentry_open+0x449/0x1000
[ 66.646132] ? inode_permission+0x125/0x560
[ 66.646148] ? cdev_device_add+0x1f0/0x1f0
[ 66.646167] vfs_open+0x9f/0xd0
[ 66.646181] path_openat+0xd58/0x3eb0
[ 66.646199] ? kasan_save_alloc_info+0x1e/0x30
[ 66.646212] ? _kasan_slab_alloc+0x90/0xa0
[ 66.646229] ? kmem_cache_alloc+0x1f6/0x310
[ 66.646244] ? getname_flags.part.0+0x52/0x490
[ 66.646260] ? getname+0x7a/0xb0
[ 66.646275] ? _x64_sys_openat+0x128/0x210
[ 66.646288] ? do_syscall_64+0x59/0x90
[ 66.646306] ? path_lookupat+0x660/0x660
[ 66.646328] do_filp_open+0x1b1/0x3e0
[ 66.646344] ? debug_smp_processor_id+0x17/0x20
[ 66.646363] ? may_open_dev+0xd0/0xd0
[ 66.646378] ? getname_flags.part.0+0x52/0x490
[ 66.646399] ? _raw_spin_lock_bh+0xe0/0xe0
[ 66.646422] ? do_sys_openat2+0x132/0x450
[ 66.646435] ? recalc_sigpending+0x144/0x1c0
[ 66.646450] ? build_open_flags+0x450/0x450
[ 66.646465] ? sigprocmask+0x201/0x360
[ 66.646478] ? __ia32_sys_ssetmask+0x1d0/0x1d0
[ 66.646495] ? _x64_sys_openat+0x128/0x210
[ 66.646509] ? __ia32_compat_sys_open+0x1b0/0x1b0
[ 66.646526] ? debug_smp_processor_id+0x17/0x20
[ 66.646545] do_syscall_64+0x59/0x90
[ 66.646560] ? syscall_exit_to_user_mode+0x26/0x50
[ 66.646578] ? do_syscall_64+0x69/0x90
[ 66.646593] ? debug_smp_processor_id+0x17/0x20
[ 66.646610] ? fpregs_assert_state_consistent+0x52/0xc0
[ 66.646630] ? exit_to_user_mode_prepare+0x49/0x1a0
[ 66.646644] ? syscall_exit_to_user_mode+0x26/0x50
[ 66.646662] ? do_syscall_64+0x69/0x90
[ 66.646675] ? exit_to_user_mode_prepare+0x16a/0x1a0
[ 66.646689] ? syscall_exit_to_user_mode+0x26/0x50
[ 66.646706] entry_SYSCALL_64_after_hwframe+0x63/0xcd
[ 66.646722] RIP: 0033:0x4534e4
[ 66.646736] Code: 24 20 eb 8f 66 90 44 89 54 24 0c e8 d6 ab 02 00 44 8b 54 24 0c 44 89 e2 48 89 ee 41 89 c0 bf 9c ff ff ff b8 01 01 00 00 0f 05 <48> 3d 00 f0 ff ff 77 34 44 89 c7 89 44
24 0c e8 18 ac 02 00 8b 44
[ 66.646751] RSP: 002b:00007f8130000140 EFLAGS: 00000293 ORIG_RAX: 0000000000000101
[ 66.646771] RAX: ffffffff8130000640 RBX: 00007f8130000640 RCX: 000000000004534e4
[ 66.646784] RDX: 0000000000000000 RSI: 00000000004b1008 RDI: 00000000ffffff9c
[ 66.646795] RBP: 00000000004b1008 R08: 0000000000000000 R09: 00007f822c66baf
[ 66.646806] R10: 0000000000000000 R11: 0000000000000293 R12: 0000000000000000
[ 66.646817] R13: 0000000000000000 R14: 000000000041b2a0 R15: 00007f812f800000
[ 66.646834] </TASK>

[ 66.646848] Allocated by task 2115:
[ 66.646859] kasan_save_stack+0x26/0x50
[ 66.646875] kasan_set_track+0x25/0x40
[ 66.646888] kasan_save_alloc_info+0x1e/0x30
[ 66.646898] kasan_kmalloc+0xb4/0xc0
[ 66.646911] kmalloc_trace+0x4a/0xb0
[ 66.646924] xillybus_init_chrdev+0xf2/0x815 [xillybus_class]
[ 66.646944] xillyusb_probe.cold+0xa61/0xadf [xillyusb]
[ 66.646966] usb_probe_interface+0x266/0x740
[ 66.646981] really_probe+0x1fa/0xa80
[ 66.646993] driver_probe_device+0x2cb/0x490
[ 66.647005] driver_probe_device+0x4e/0x140
[ 66.647017] driver_attach+0x1a3/0x520
[ 66.647028] bus_for_each_dev+0x11e/0x1c0
[ 66.647039] driver_attach+0x3d/0x60
[ 66.647050] bus_add_driver+0x449/0x5a0
[ 66.647061] driver_register+0x219/0x390
[ 66.647073] usb_register_driver+0x228/0x400
[ 66.647084] 0xffffffffc033202d
[ 66.647094] do_one_initcall+0x97/0x310
[ 66.647109] do_init_module+0x19a/0x630
[ 66.647120] load_module+0x6ca4/0x7d90
[ 66.647131] __do_sys_finit_module+0x134/0x1d0
[ 66.647142] __x64_sys_finit_module+0x72/0xb0
[ 66.647153] do_syscall_64+0x59/0x90
[ 66.647164] entry_SYSCALL_64_after_hwframe+0x63/0xcd

[ 66.647182] Freed by task 2089:
[ 66.647191] kasan_save_stack+0x26/0x50
[ 66.647205] kasan_set_track+0x25/0x40
[ 66.647218] kasan_save_free_info+0x2e/0x50
[ 66.647228] kasan_slab_free+0x174/0x1e0
[ 66.647241] kasan_slab_free+0x12/0x20
[ 66.647255] slab_free_freelist_hook+0xd0/0x1a0
[ 66.647267] kmem_cache_free+0x193/0x2c0
[ 66.647280] kfree+0x29/0x120
[ 66.647291] xillybus_cleanup_chrdev+0x3c1/0x570 [xillybus_class]
[ 66.647311] xillyusb_disconnect+0xfe/0x790 [xillyusb]
[ 66.647332] usb_unbind_interface+0x187/0x7c0
[ 66.647345] device_remove+0x117/0x170
[ 66.647357] device_release_driver_internal+0x418/0x660
[ 66.647369] device_release_driver+0x12/0x20
[ 66.647381] bus_remove_device+0x28f/0x540
[ 66.647392] device_del+0x501/0xc30
[ 66.647407] usb_disable_device+0x2a5/0x660
[ 66.647418] usb_disconnect.cold+0x1f9/0x620
[ 66.647431] hub_event+0x16d3/0x3d20
[ 66.647446] process_one_work+0x778/0x11c0
[ 66.647459] worker_thread+0x544/0x1180
[ 66.647471] kthread+0x280/0x320
[ 66.647481] ret_from_fork+0x1f/0x30

[ 66.647501] Last potentially related work creation:
[ 66.647510] kasan_save_stack+0x26/0x50
[ 66.647524] kasan_record_aux_stack+0xb6/0xd0
[ 66.647534] kasan_record_aux_stack_noalloc+0xb/0x20
[ 66.647544] kvfree_call_rcu+0xaf/0xa50
[ 66.647555] memcg_reparent_list_lrus+0x477/0x790
[ 66.647566] mem_cgroup_css_offline+0x1ea/0x310
[ 66.647579] css_killed_work_fn+0x4/0x30
[ 66.647590] process_one_work+0x778/0x11c0
[ 66.647602] worker_thread+0x544/0x1180
```

```
[ 66.647614] kthread+0x280/0x320
[ 66.647623] ret_from_fork+0x1f/0x30

[ 66.647642] The buggy address belongs to the object at ffff88811eb45f80
           which belongs to the cache kmalloc-64 of size 64
[ 66.647656] The buggy address is located 16 bytes inside of
           64-byte region [ffff88811eb45f80, ffff88811eb45fc0)

[ 66.647677] The buggy address belongs to the physical page:
[ 66.647686] page:000000009bdbcfb refcount:1 mapcount:0 mapping:0000000000000000 index:0xffff88811eb45a00 pfn:0x11eb45
[ 66.647701] flags: 0x17ffffc0000200(slab|node=0|zone=2|lastcpupid=0x1fffff)
[ 66.647720] raw: 0017ffffc0000200 ffffea000481fec8 ffffea0004544b48 ffff888100042640
[ 66.647731] raw: ffff88811eb45a00 000000000020001e 00000001fffffff 0000000000000000
[ 66.647738] page dumped because: kasan: bad access detected


[ 66.647751] Memory state around the buggy address:
[ 66.647760] ffff88811eb45e80: 00 00 00 00 00 00 00 00 fc fc fc fc fc fc fc fc
[ 66.647771] ffff88811eb45f00: 00 00 00 00 00 fc fc fc fc fc fc fc fc fc fc
[ 66.647781] >ffff88811eb45f80: fa fb fb fb fb fb fb fb fc fc fc fc fc fc fc
[ 66.647790]                                     ^
[ 66.647800] ffff88811eb46000: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
[ 66.647810] ffff88811eb46080: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
[ 66.647818] =====
[ 66.647826] Disabling lock debugging due to kernel taint
...\n
```

Regards,
Hyunwoo Kim.

[next](#) [reply](#) other threads:[~2022-10-22 17:54 UTC|newest]

Thread overview: 4+ messages / [expand](#)[flat|nested] [mbox.gz](#) [Atom](#) [feed](#) [top](#)
2022-10-22 17:54 Hyunwoo Kim [this message]
2022-10-23 14:19 [\[PATCH\]](#) char: xillybus: Fix use-after-free in xillyusb_open() Eli Billauer
2022-10-23 14:26 Hyunwoo Kim
2022-10-24 7:10 Eli Billauer

find likely ancestor, descendant, or conflicting patches for [this message](#):

[dfblob:39bcbfd908b](#) [dfblob:d615f74dcf3](#) 

 [\(help\)](#)

Reply instructions:

You may reply publicly to [this message](#) via plain-text email
using any one of the following methods:

* Save the following mbox file, import it into your mail client,
and reply-to-all from there: [mbox](#)

Avoid top-posting and favor interleaved quoting:
https://en.wikipedia.org/wiki/Posting_style#Interleaved_style

* Reply using the **--to**, **--cc**, and **--in-reply-to**
switches of `git-send-email(1)`:

```
git send-email \
  --in-reply-to=20221022175404.GA375335@ubuntu \
  --to=inv4bel@gmail.com \
  --cc=arnd@arndb.de \
  --cc=eli.billauer@gmail.com \
  --cc=gregkh@linuxfoundation.org \
  --cc=linux-kernel@vger.kernel.org \
  /path/to/YOUR_REPLY
```

<https://kernel.org/pub/software/scm/git/docs/git-send-email.html>

* If your mail client supports setting the **In-Reply-To** header
via mailto: links, try the [mailto: link](#)

Be sure your reply has a **Subject:** header at the top and a blank line before the message body.

This is an external index of several public inboxes,
see [mirroring instructions](#) on how to clone and mirror
all data and code used by this external index.