

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow @Openwall on Twitter for new release announcements and other news

[<prev](#) [\[next>\]](#) [<thread-prev](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Tue, 29 Nov 2022 21:59:55 +0100  
From: Julien Pivotto <roidelapluie@...metheus.io>  
To: Solar Designer <solar@...nwall.com>  
Cc: oss-security@...ts.openwall.com  
Subject: Re: CVE-2022-46146 in Prometheus' exporter toolkit:  
bypass basic authentication

On 29 Nov 16:38, Solar Designer wrote:  
> On Tue, Nov 29, 2022 at 01:22:58PM +0100, Julien Pivotto wrote:  
> > The exporter toolkit is a go library intended at Prometheus exporters.  
> > It provides some features that are useful for Prometheus exporters,  
> > which work by exposing HTTP servers to be exposed by the Prometheus  
> > server.  
> >  
> > One of those features is basic authentication. To achieve this,  
> > Prometheus requires you to store a bcrypt hash into a file, web.yml.  
> >  
> > While bcrypt is fine, it takes by design a lot of time and resources to  
> > compare a password with a hash. To limit this impact, we have a built-in  
> > cache that caches the good and bad answers.  
> >  
> > Once a request comes, we check it against the cache and decide whether  
> > to allow the request. We also check that the user is valid. However, the  
> > key for that cache is predictable:  
> >  
> > hex(username + hashed password + input password)  
> >  
> > If you know the bcrypted password, you can poison the cache and use that  
> > cached positive value in a subsequent query:  
> >  
> > Request 1:  
> >  
> > username = username+hashed password  
> > password = "fakepassword"  
> >  
> > Request 2:  
> >  
> > username = username  
> > password = bcrypt(fakepassword)+"fakepassword"  
> >  
> > "fakepassword" is used as bcrypted password when a user does not exist.  
> >  
> > The fact that we save unhappy tentatives and that we validate  
> > non-existing users against "fakepassword" is to prevent side channel  
> > attacks that could reveal if a user exists in a system or not.  
> >  
> > Prometheus 2.37.4 and 2.40.4 are out, with this fix. We recommend all  
> > the exporters that depend on the repository to upgrade.  
> >  
> > CVE-2022-46146 was assigned to this security report in our exporter  
> > toolkit:  
> > <https://github.com/prometheus/exporter-toolkit/security/advisories/GHSA-7rg2-cxvp-9p7p>  
> >  
> > We would like to thank Lei Wan for the responsible disclosure of this  
> > bug.  
> >  
> > The above describes the issue, but not the fix. This left me curious.  
> >  
> > The fix commit appears to be this:  
> >  
> > <https://github.com/prometheus/exporter-toolkit/commit/5bleab34484ddd353986bce736cd119d863e4ff5>  
> >  
> > It makes two changes:  
> >  
> > 1. Rather than concatenate the original strings to produce the cache  
> > key, the 3 individual components are first hex-encoded and are then  
> > concatenated with colons as separators.  
> >  
> > 2. Cache records for authentication against non-existent users with  
> > "fakepassword" no longer indicate that authentication passed.  
> >  
> > This appears sufficient to address the described issue.  
> >  
> > The caching is controversial. A comment in cache.go says:  
> >  
> > // newCache returns a cache that contains a mapping of plaintext passwords  
> > // to their hashes (with random eviction). This can greatly improve the  
> > // performance of traffic-heavy servers that use secure password hashing  
> > // algorithms, with the downside that plaintext passwords will be stored in  
> > // memory for a longer time (this should not be a problem as long as your  
> > // machine is not compromised, at which point all bets are off, since basicauth  
> > // necessitates plaintext passwords being received over the wire anyway).  
> >  
> > IMO, storage of plaintext passwords in memory for longer doesn't become  
> > a non-issue just because plaintext passwords are also available during  
> > authentication. A compromise might be short-lived and not every user  
> > who had logged in before would necessarily log in again while the server  
> > is compromised, so storage of plaintext passwords that might have been  
> > used a long time ago does make things worse and partially defeats the  
> > purpose of password hashing. OTOH, in a persistent Go service they're  
> > likely to stay around anyway.  
> >  
> > Another (minor) concern is that the cache is indexed by password-derived  
> > material, making the password a bit more susceptible to local CPU cache  
> > timing attacks (after a leak of bcrypt's random salts to the attacker).  
> > bcrypt itself also does such indexing, but that's part of why it turned  
> > out to be relatively inefficient on GPUs, so it can be viewed as  
> > justified risk. Is the risk from the cache also justified, by it saving  
> > computing resources (not under deliberate DoS, though)? Maybe.  
> > Alternative designs of the cache are possible, but involve other  
> > non-trivial trade-offs and subtle detail. If Go uses keyed hashing for  
> > its maps (does it? I don't know), that also provides a mitigation (while  
> > the random key is not leaked/inferred).  
> >  
> > Further, perhaps Go maps internally compare the provided key against  
> > some stored keys (within one hash bucket?), and that can probably  
> > involve a byte-by-byte comparison, kind of leaking the length of matched  
> > substring even to the remote client and potentially allowing to probe  
> > candidate passwords character-by-character (when hitting the same hash  
> > bucket). This is mitigated by such probing also thrashing the cache and  
> > triggering the much slower bcrypt computation.  
> >  
> > To clarify, I am not suggesting that any changes to the code be made -  
> > they could as well introduce new issues, and would need a new review.  
> > Dropping of the cache is a pretty obvious change to make, but maybe the  
> > risks involved are no big deal for this specific service.

Thanks for your meaningful message.

I want to clarify the use case of this library for users who are not familiar with Prometheus. The main goal is to secure communication between two components (Prometheus and exporters) that speak every 10-30s, generally with a single (machine) user.

Another goal is to secure the Prometheus server itself - where users are actually connecting. But in general, you would use a proper proxy in front to deal with more secure protocols than basic auth - such as OpenID connect - and those proxies would communicate with Prometheus either via TLS+Basic auth or TLS+Client-certificate or via loopback interface.

On my machine, in a quick test, the first request (using Node Exporter) would take 103ms. Thanks to the cache, subsequent requests take ~58ms. The cache produces a significant improvement with resource usage and time.

What I have in mind is that we could make the cache optional - but outside of the scope of this pull request; this would be a new feature.

>  
> I was just curious and I thought some others in here would be as well,  
> so I am sharing observations and thoughts.  
>  
> Alexander

--  
Julien Pivotto  
@roidelapluie

**Download attachment "[signature.asc](#)" of type "application/pgp-signature" (229 bytes)**

[Powered by blists](#) - [more mailing lists](#)

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

